

Agents that prove their numbers are right.

Numbers from code. A separate auditor re-runs the math, catching a report that is wrong *even when it looks honest.*

numbers from code · AI directs · code computes · humans decide

- Band-native
- 3 agents
- deterministic engine
- correctness audit

Everyone audits *honesty*.
We prove **correctness**.

EVERYONE ELSE Honesty checks and tamper-evident hashes. Proves nobody <i>altered</i> it.	ASSAY Re-runs the engine, diffs every figure. Proves it is <i>right</i>, not just honest.
--	---

A hash proves nobody altered it. **Not** that it is right.

Three agents collaborate through Band

Skaut scout · discovery	Konduktor orchestrate · write	engine ratios · verdict	Pengulas re-run · audit
---	---	---	---

✓ AUDIT: PASS

ORKES-A · B · C · E

Checks 7 | Discrepancies 0

X AUDIT: FAIL

ORKES-A · Discrepancies 1

Details: gearing: report claims 0.01, engine computes 0.26

The verdict lives in **code**, not the model

No model touches a number, so the verdict is byte-identical across brains.

```
MODEL_PROVIDER = glm | claude | featherless | aimlapi    · one env var, no code edit

GLM-5.2 (z.ai)      ✓ validated live · 4 PASS · 2 refused
Claude (Opus 4.8)  ✓ validated live · 4 PASS · 2 refused
Qwen3 · GPT-4o     same verdict by construction
```

- ❑ Refused **2 of 6** as INCOMPLETE. No data, no number, no report.
- ❑ Screens and audits only. Ends AWAITING HUMAN SIGN-OFF. Manusia memutuskan.

The *wrong-but-honest* number is the expensive one

A wrong-but-honest figure clears review. A restatement then costs about **9% of share price**. (Palmrose 2004)

The buyer

Compliance, risk and internal-audit desks. They sign off, and own the consequence.

Horizontal

Only the engine is domain-specific. Same control: underwriting, claims, ESG, tax.

24 tests green • deterministic auditor • model-portable • synthetic demo data • not investment advice