Agents that prove their numbers are right.

Numbers from code. A separate auditor re-runs the math, catching a report that is wrong *even when it looks honest.*

numbers from code · AI directs · code computes · humans decide

Band-native | 3 agents | deterministic engine | correctness audit

Everyone audits *honesty*.
We prove correctness.

EVERYONE ELSE	ASSAY
Honesty checks and tamper-evident hashes. Proves nobody *altered* it.	Re-runs the engine, diffs every figure. Proves it is *right*, not just honest.

A hash proves nobody altered it. **Not** that it is right.

Three agents collaborate through Band

Skaut	Konduktor	engine	Pengulas
scout · discovery	orchestrate · write	ratios · verdict	re-run · audit

✓ **AUDIT: PASS** ORKES-A · B · C · E Checks 7 | Discrepancies 0

✗ **AUDIT: FAIL** ORKES-A · Discrepancies 1
Details: gearing: report claims 0.01, engine computes 0.26

The verdict lives in code, not the model

No model touches a number, so the verdict is byte-identical across brains.

```
MODEL_PROVIDER = claude | glm    · one env var, no code edit

Claude (Opus 4.8)   ✓ validated live · 4 PASS · 2 refused
GLM-5.2 (z.ai)      ✓ validated live · 4 PASS · 2 refused
Any other brain     same verdict by construction
```

- Refused 2 of 6 as INCOMPLETE. No data, no number, no report.
- Screens and audits only. Ends AWAITING HUMAN SIGN-OFF. Manusia memutuskan.

The *wrong-but-honest* number is the expensive one

A wrong-but-honest figure clears review. A restatement then costs about **9% of share price.** (Palmrose 2004)

The buyer

Compliance, risk and internal-audit desks. They sign off, and own the consequence.

Horizontal

Only the engine is domain-specific. Same control: underwriting, claims, ESG, tax.

24 tests green · deterministic auditor · model-portable · synthetic demo data · not investment advice