

Band Armada

Agents that **refuse to invent numbers**.

Numbers from code. AI directs, code computes, humans decide.

Band-native

3 agents

deterministic engine

correctness audit

Most agents happily invent a number

The smarter ones add an auditor that checks whether the AI was honest about its number.

That still passes a verdict that is **wrong** as long as it looks honest. In a high-stakes financial workflow that is a liability, not a quirk.

We compute the decision in **code**

So our auditor re-runs the math and catches a verdict that is wrong **even when the AI is perfectly honest.**


Everyone else

Honesty / citation / tamper-evident hash. Proves nobody altered the report.

Band Armada

Correctness audit. Re-computes the number and proves the report is right.

Three agents collaborate **through Band**



- No agent computes a ratio. The engine is the only source of numbers.
- Every @mention, tool_call and tool_result is in Band's audit trail.
- Incomplete data is refused: no number, no report.

Clean run, then the tamper catch

4 reports, audited by Pengulas re-running the engine:

```
AUDIT: PASS  ORKES-A · ORKES-B · ORKES-C · ORKES-E  (Checks 7 |  
Discrepancies 0)
```

Doctor one figure in a finished report:

```
AUDIT: FAIL  Ticker: ORKES-A · Discrepancies: 1  
Details: gearing: report claims 0.01, engine computes 0.26
```

Refusal-first & human sign-off

- Refused 2 of 6 companies (TIDAK LENGKAP) on incomplete data. Most desks would score them anyway.
- The desk screens and audits only. It does not buy, sell, or act.
- Every round ends **AWAITING HUMAN SIGN-OFF** — manusia memutuskan.
- Brain is swappable by one env var. The decision does not live in the model.

Agents that prove their numbers are right.

Band · LangGraph

deterministic engine

24 tests green

model-portable

Built zero-extra-cost on synthetic demo data. Not investment advice.